



UNIVERSITY OF
ARKANSAS

Convolutional Neural Networks

Jiahui Chen

Department of Mathematical Sciences

University of Arkansas

Reference: Ming Li's notes

Introduction


- Fukushima (1980) – Neo-Cognitron; LeCun (1998) – Convolutional Neural Networks (CNN, or ConvNet);...

Motivation:

Images typically have 1000^2 pixels, which give rise to 1000^2 data points or features, leading to an intractably high dimension of the weight space. However, not all of them are essential due to the spatial patterns. Many of them have shared properties. Therefore, the weight space dimension can be dramatically reduced if an appropriate pre-processing of the image data is carried out to analyze or extract spatial correlations or patterns in images.

Motivation

- We know it is good to learn a small model.
- From this fully connected model, do we really need all the edges?
- Can some of these be shared?




Motivation

Consider learning an image:



Some patterns are much smaller than the whole image

Can represent a small region with fewer parameters




Same pattern appears in different places:
They can be compressed!

What about training a lot of such “small” detectors and each detector must
“move around”.



A Convolutional Layer

A CNN is a neural network with some convolutional layers (and some other layers). A convolutional layer has a number of filters that does convolutional operation.



Convolution

These are the network parameters to be learned.

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

⋮ ⋮

Each filter detects a small pattern (3 x 3).

Convolution

stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

Dot
product



3

-1

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

Convolution

If stride=2

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

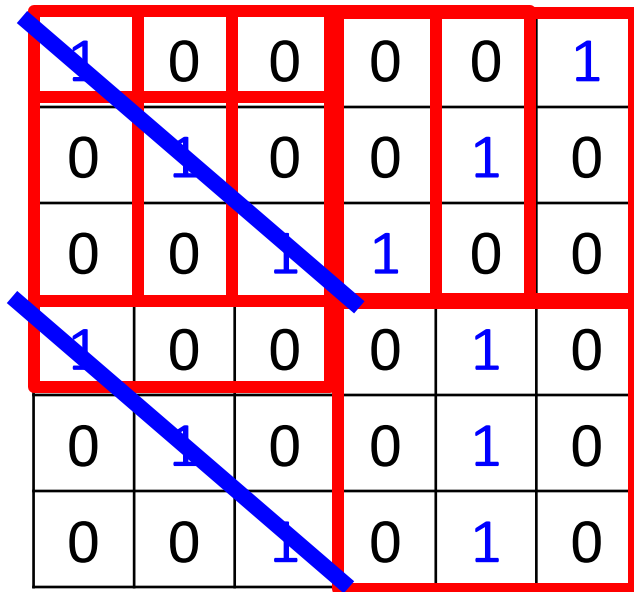
6 x 6 image

3 -3

1	-1	-1
-1	1	-1
-1	-1	1

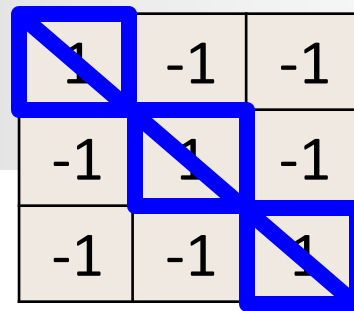
Filter 1

stride=1



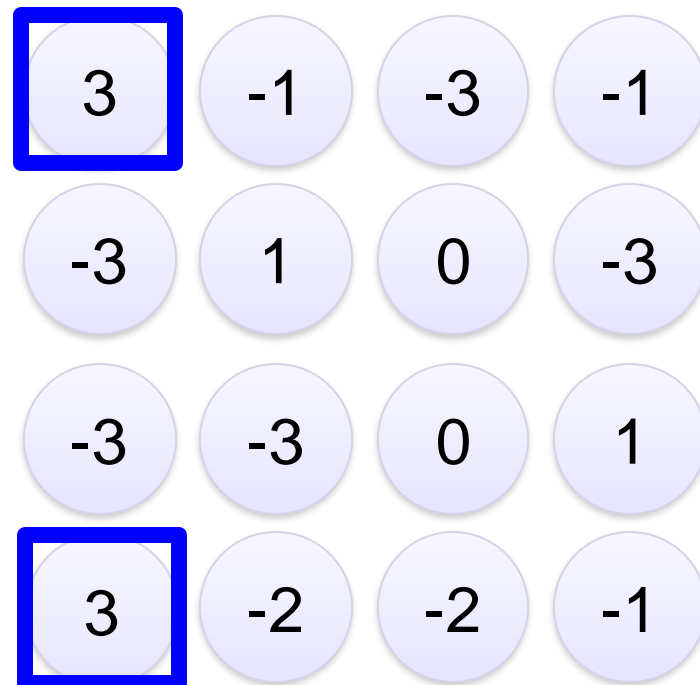
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image



1	-1	-1
-1	1	-1
-1	-1	1

Filter 1



3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

stride=1


1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

Repeat this for each filter



Two 4 x 4 images Forming 2 x 4 x 4 matrix


The diagram illustrates three overlapping 3x3 grids. The top grid is mostly empty. The middle grid contains the following values:

1	-1	-1
-1	1	-1
-1	-1	1

The bottom grid is mostly empty.

Filter 1

Filter 2



The diagram shows a 6x6 grid of cells. The top row of the grid contains the numbers 1, 0, 0, 0, 0, 1 in blue. The other rows contain 0s and 1s in blue. To the left of the grid, there is a vertical sequence of numbers: 1, 0, 0, 0, 0, 1, which are aligned with the rows of the grid. The numbers 1 and 0 are in blue, while the 0s in the other rows are in black.

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

image

1	-1	-1
-1	1	-1
-1	-1	1

-1	1	-1
-1	1	-1
-1	1	-1





convolution

-1	-1	-1	-1
-1	-1	-2	1
-1	-1	-2	1
-1	0	-4	3


Fully-
connected


1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0





fewer parameters!






6 x 6 image

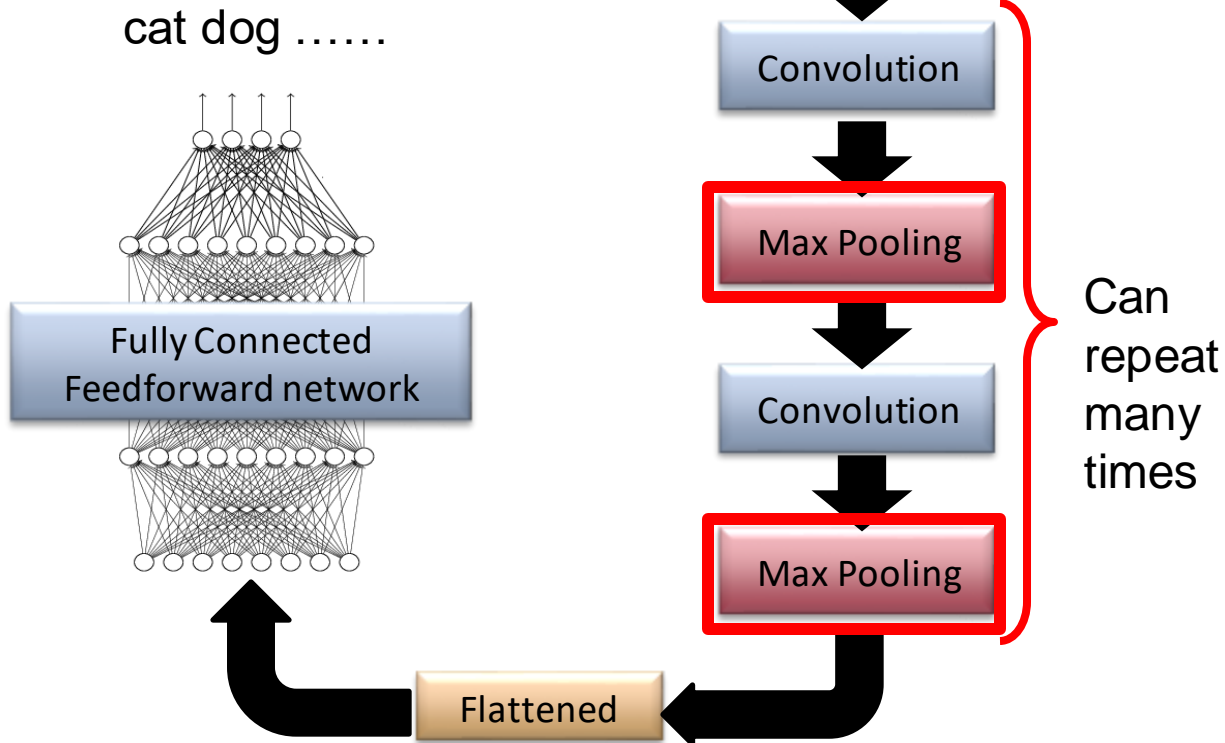
Fewer parameters

Even fewer parameters



Shared weights

The whole CNN



Max Pooling

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

3	-	-	-
1	3	1	-
-	0	-	3
3	1		
-	3	0	1
3	3	-	-
3	-	2	1
2	2	-	-

-	-	-	-
1	1	1	1
-	-	-	1
1	1	2	
-	-	-	1
1	1	2	1
-	0	-	3
1	4	-	-

Why Pooling

- Subsampling pixels will not change the object

bird



Subsampling

bird



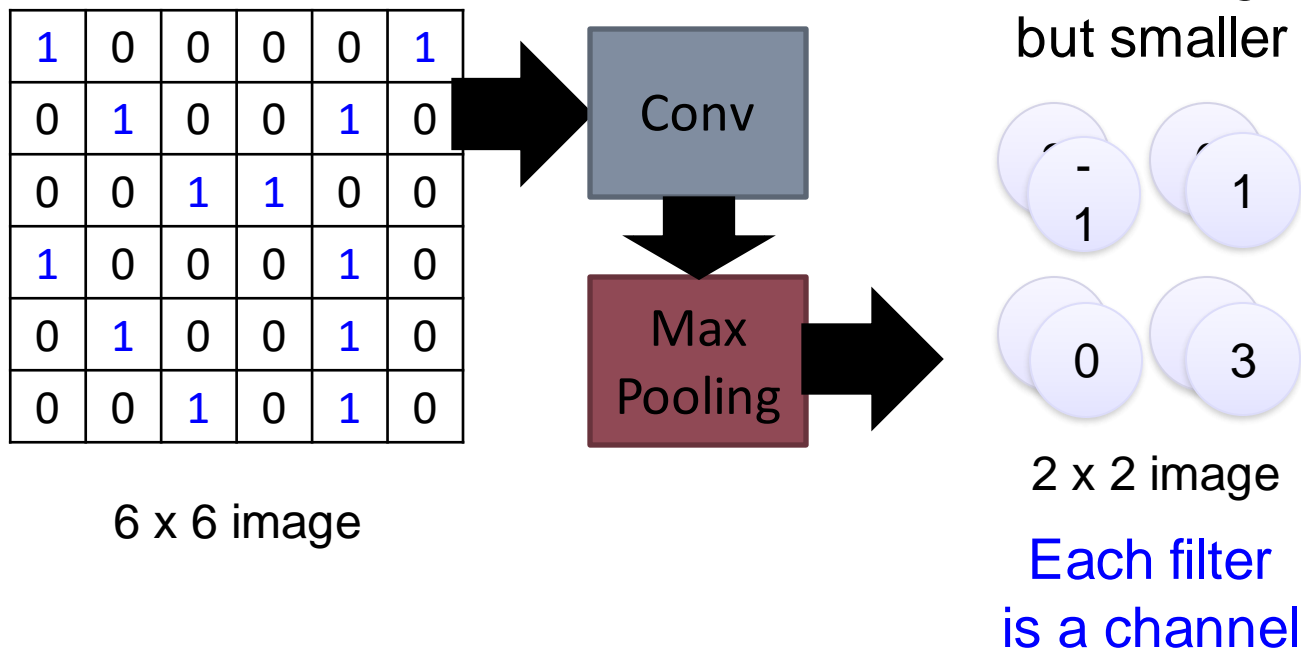
We can subsample the pixels to make image
smaller → fewer parameters to characterize the image




A CNN compresses a fully connected network in two ways:

- Reducing number of connections
- Shared weights on the edges
- Max pooling further reduces the complexity

Max Pooling




The whole CNN



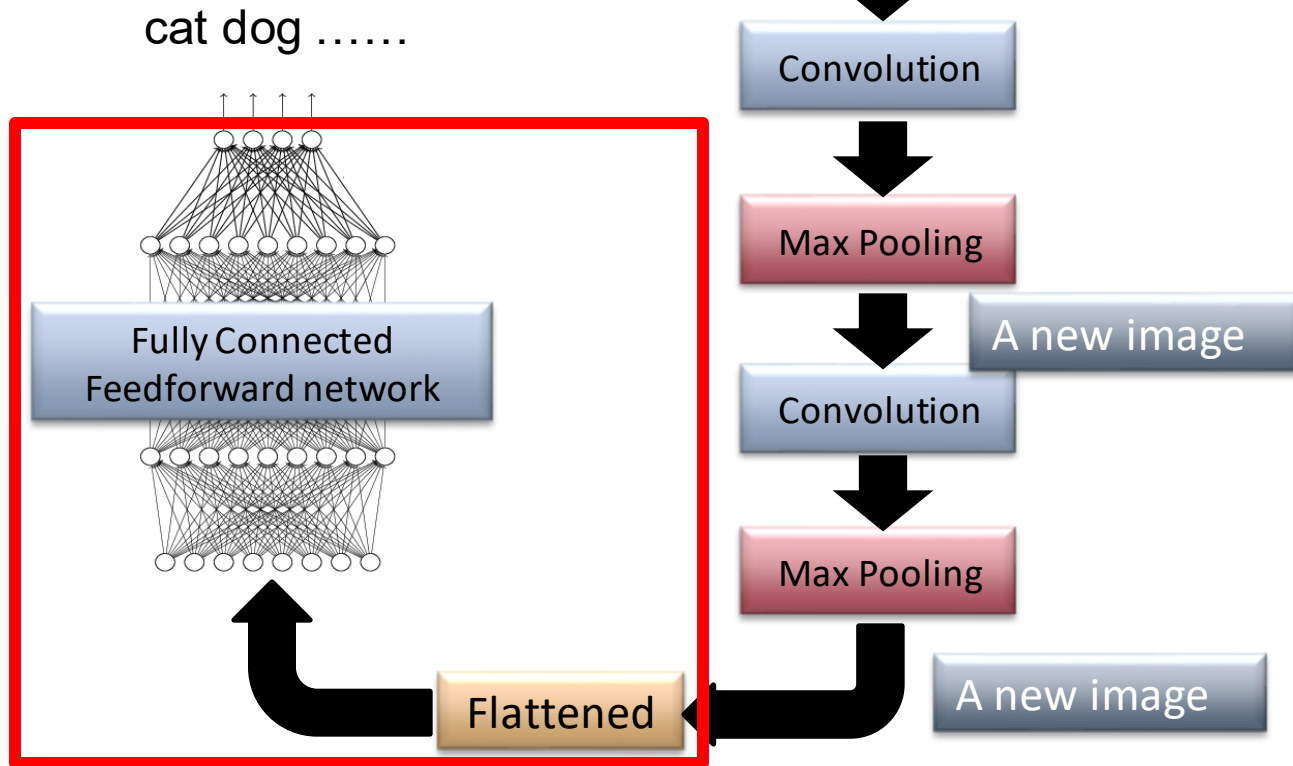
Smaller than the original image

The number of channels is the number of filters

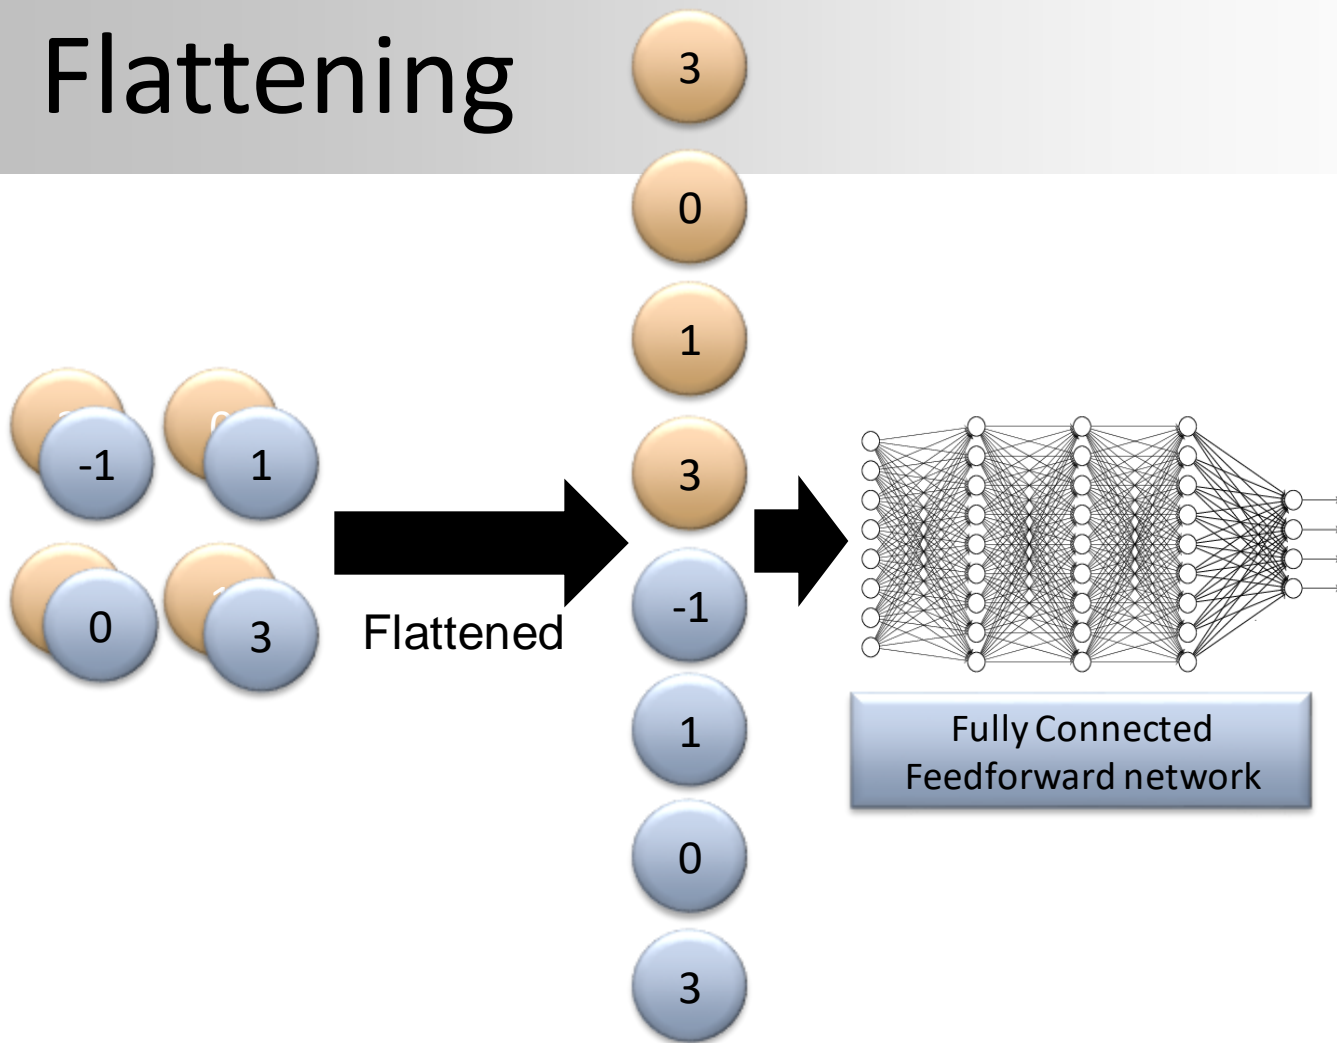


Can repeat many times

The whole CNN



Flattening



CNN in Keras

Only modified the **network structure** and **input format (vector -> 3-D tensor)**

```
model2.add( Convolution2D( 25, 3, 3,
                           input_shape=(28, 28, 1)) )
```

1	-1	-1	1	-1
-1	1	-1	1	-1
-1	-1	-1	1	-1




There are
25 3x3
filters.

Input_shape = (28 , 28 , 1)

28 x 28 pixels

1: black/white, 3: RGB

```
model2.add( MaxPooling2D( (2, 2) ) )
```



input



Convolution



Max Pooling



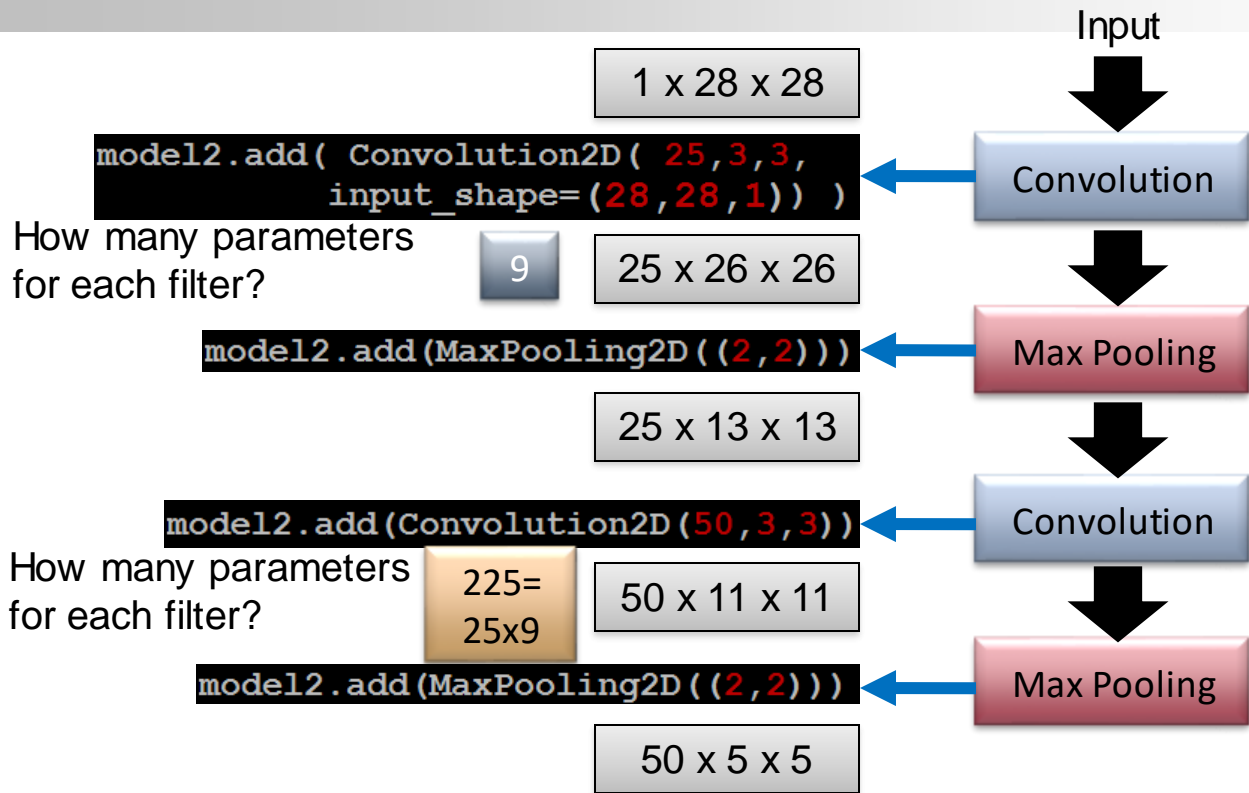
Convolution



Max Pooling

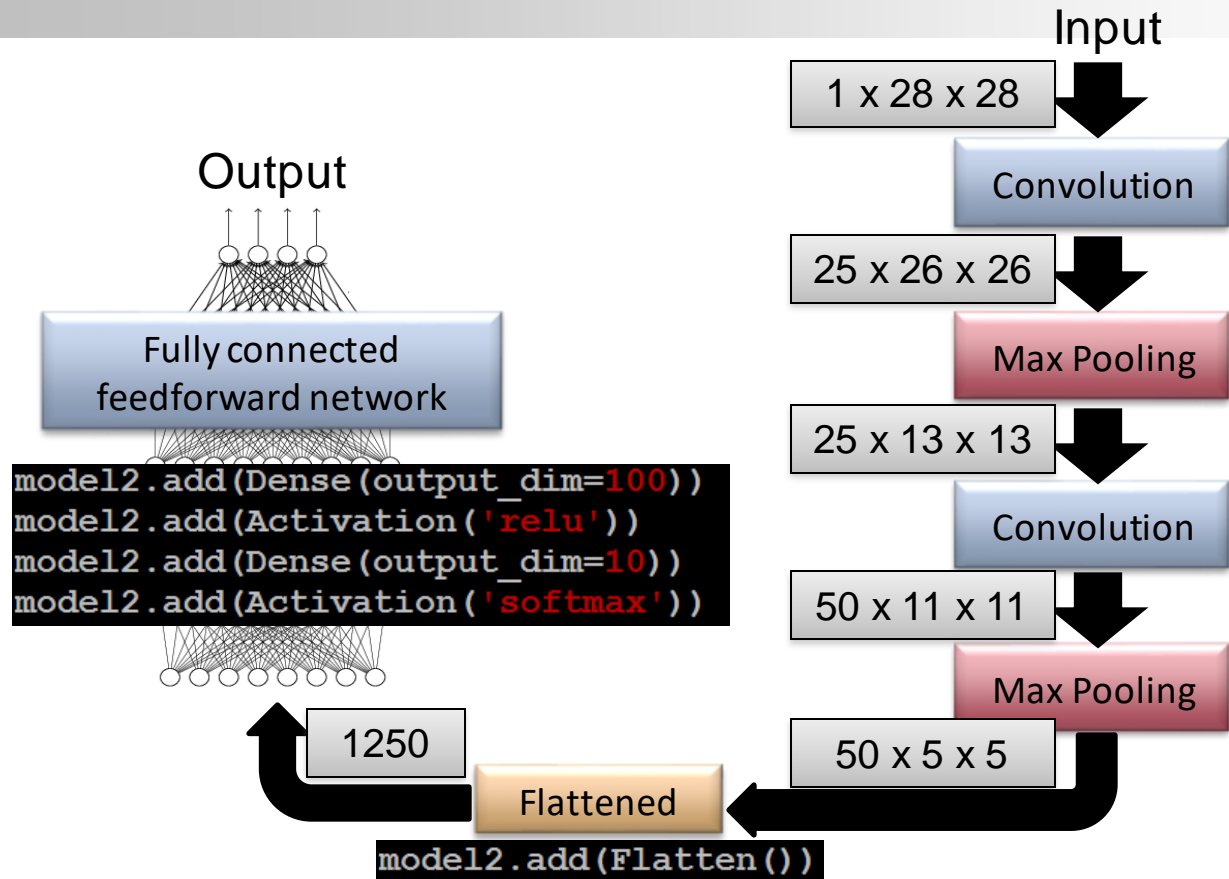
CNN in Keras

Only modified the *network structure* and *input format* (vector -> 3-D array)



CNN in Keras

Only modified the *network structure* and *input format (vector -> 3-D array)*



AlphaGo




19 x 19 matrix

Black: 1

white: -1

none: 0



Next move
(19 x 19
positions)

Fully-connected feedforward network
can be used

But CNN performs much better

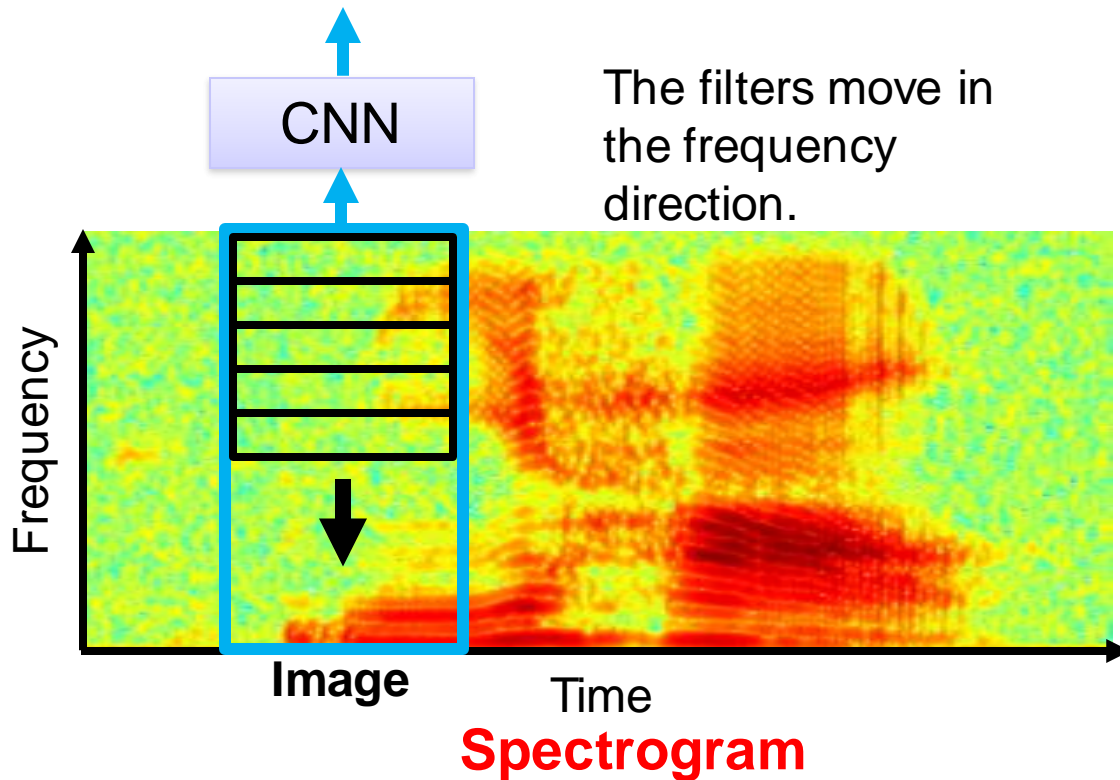
AlphaGo's policy network

The following is quotation from their Nature article:

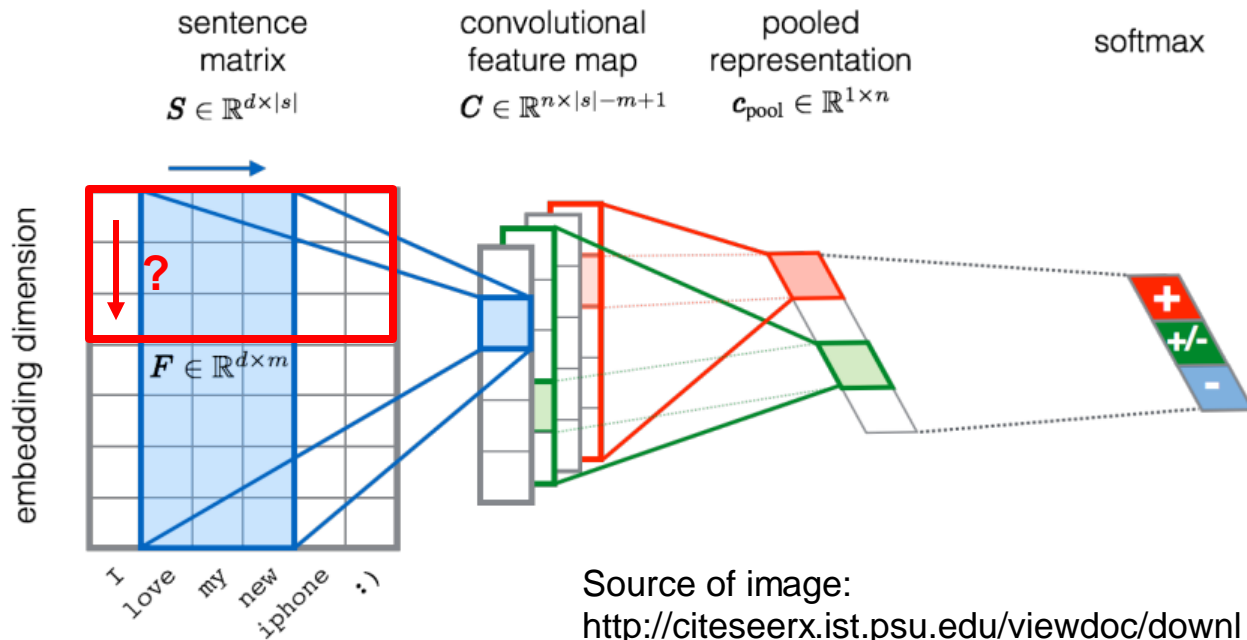
Note: AlphaGo does not use Max Pooling.

Neural network architecture. The input to the policy network is a $19 \times 19 \times 48$ image stack consisting of 48 feature planes. The first hidden layer zero pads the input into a 23×23 image, then convolves k filters of kernel size 5×5 with stride 1 with the input image and applies a rectifier nonlinearity. Each of the subsequent hidden layers 2 to 12 zero pads the respective previous hidden layer into a 21×21 image, then convolves k filters of kernel size 3×3 with stride 1, again followed by a rectifier nonlinearity. The final layer convolves 1 filter of kernel size 1×1 with stride 1, with a different bias for each position, and applies a softmax function. The match version of AlphaGo used $k = 192$ filters; Fig. 2b and Extended Data Table 3 additionally show the results of training with $k = 128, 256$ and 384 filters.

CNN in speech recognition



CNN in text classification



Source of image:
<http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.703.6858&rep=rep1&type=pdf>